

# Migration Plan: Crema.co (Heroku to AWS)

Client: Crema.co

Document Ref: MIGRATION-PLAN-V1

## 1. Executive Summary

This document outlines the technical execution plan for migrating the Crema.co Ruby on Rails e-commerce platform from Heroku to Amazon Web Services (AWS). The objective is to transition the monolithic application to a scalable, containerized architecture using AWS ECS Fargate, Amazon RDS, and ElastiCache, while ensuring minimal downtime and robust security via private networking (NAT Gateways) and Application Load Balancers (ALB).

## 2. Prerequisites & Access Verification

Before initiating the migration, the following access points have been confirmed:

- **Heroku:** Full access to the existing production application and configuration.
- **GitHub:** Access to the master branch containing the source code and Dockerfile.
- **AWS:** Administrator access to the target AWS account for resource provisioning.

## 3. Target Architecture Overview

Component	Strategy	Details
Compute	ECS Fargate	Serverless container execution. Eliminates need for EC2 instance management.
Database	Amazon RDS	PostgreSQL instance (provisioned to match or exceed Heroku spec).
Caching	ElastiCache	Redis Cluster mode enabled (if required) or Primary/Replica setup.
Networking	VPC with NAT	Public Subnets: ALB, NAT

		Gateways. <b>Private Subnets:</b> ECS Tasks, RDS, ElastiCache.
<b>Load Balancing</b>	<b>ALB</b>	Application Load Balancer handling SSL termination and traffic distribution.

## 4. Technical Migration Phases

### Phase 1: Application Containerization & Configuration

**Objective:** Finalize the Docker image and secure application configuration.

1. **Dockerfile Finalization:**
  - Utilize the Dockerfile located in the master branch.
  - Apply the necessary tweaks identified during initial testing.
  - **Action:** Build and tag the image locally to verify the build process completes successfully.
2. **Environment Variable Retrieval (Critical Path):**
  - *Current Issue:* Local tests are failing due to missing environment variables.
  - **Storage:** Upload sensitive variables to **AWS Secrets Manager** or **AWS Systems Manager Parameter Store**.

### Phase 2: Infrastructure Provisioning (IaC)

**Objective:** Provision the "empty" environment ready to receive data and code.

1. **Network Setup:**
  - Create VPC.
  - Deploy **Internet Gateways** for public subnets.
  - Deploy **NAT Gateways** in public subnets to allow private ECS tasks to pull images and reach external APIs.
2. **Database & Cache:**
  - Provision **RDS for PostgreSQL** (20GB to start to accommodate the 10GB data + growth).
  - Provision **ElastiCache for Redis**.
3. **Compute Resources:**
  - Create ECS Cluster.
  - Create ECR Repository.
  - Configure Security Groups:
    - **ALB SG:** Allow 443 from 0.0.0.0/0.

- App SG: Allow traffic only from ALB SG.
- DB/Redis SG: Allow traffic only from App SG.

## Phase 3: CI/CD Pipeline Implementation

**Objective:** Automate the build and deployment process to eliminate manual errors and ensure consistent releases.

1. **Source Stage:**
  - Connect **AWS CodePipeline** to the GitHub repository.
  - Configure a webhook to trigger the pipeline on commits to the master branch.
2. **Build Stage (AWS CodeBuild):**
  - Create a buildspec.yml file (if not present) to handle the build logic.
  - **Steps:** Log in to ECR, build the Docker image, tag the image with the commit SHA, and push the image to ECR.
  - **Output:** Generate an imagedefinitions.json file for the deploy stage.
3. **Deploy Stage (ECS):**
  - Configure CodePipeline to deploy to **Amazon ECS**.
  - Use the imagedefinitions.json artifact to update the ECS Service with the new image.
  - Configure rolling updates to ensure zero downtime during deployments.

## Phase 4: Scheduled Tasks Implementation (Heroku Scheduler Replacement)

**Objective:** Migrate recurring jobs (Rake tasks) from Heroku Scheduler to AWS Native services.

**Strategy:** Use **Amazon EventBridge Scheduler** to trigger standalone Fargate tasks using the "Container Override" pattern.

### Implementation Details:

1. **Dual Task Definition Strategy:**

To optimize resources and security, we will decouple the scheduled tasks from the main service.

  - **Task Definition 1 (crema-app):** Runs long-lived Web and Sidekiq services.
  - **Task Definition 2 (crema-rake-tasks):** Runs ephemeral scheduled tasks. This allows for tighter IAM permissions (only needs permissions relevant to the specific rake tasks) and different resource sizing (often lower CPU/Memory requirements than the web app).
2. **Task Definition Configuration (crema-rake-tasks):**
  - **Network Mode:** awsvpc (Required for Fargate).
  - **Image:** Same ECR URI as the main app (latest).
  - **Command:** Set a default placeholder command (e.g., ["echo", "EventBridge"])

Override Placeholder"]). This safeguards against accidental execution without overrides.

- **Roles:**
  - executionRoleArn: For pulling images and fetching Secrets (SSM/Secrets Manager).
  - taskRoleArn: Permissions required by the Rails app itself (e.g., S3 access).
- **Configuration Reference (JSON):**

```
{  
  "family": "crema-rake-tasks",  
  "networkMode": "awsvpc",  
  "requiresCompatibilities": ["FARGATE"],  
  "cpu": "256", "memory": "512",  
  "containerDefinitions": [  
    {  
      "name": "rake-runner",  
      "image": "<YOUR_ECR_URI>:latest",  
      "command": ["echo", "Command overridden by EventBridge"],  
      "secrets": [  
        {"name": "DATABASE_URL", "valueFrom": "arn:aws:ssm..."},  
        {"name": "REDIS_URL", "valueFrom": "arn:aws:ssm..."}  
      ],  
      "logConfiguration": {  
        "logDriver": "awslogs",  
        "options": {  
          "awslogs-group": "/ecs/crema-rake-tasks",  
          "awslogs-stream-prefix": "rake"  
        }  
      }  
    }  
  ]  
}
```

### 3. EventBridge Rules Setup:

- **Source:** Audit Heroku Scheduler for active jobs (e.g., rake daily:report, rake subscription:renew).
- **Schedule:** Create EventBridge rules using Cron expressions matching the Heroku frequency.
- **Target:** ECS Cluster -> Task Definition: crema-rake-tasks.
- **Container Overrides:** For each rule, configure the "Container Overrides" to inject the specific Rake command.

- Command: ["bundle", "exec", "rake", "subscription:renew"]
4. **End-to-End Flow Example (Daily Renewals):**
- **Trigger:** At 09:00 UTC, the EventBridge Rule crema-daily-renewals activates.
  - **Dispatch:** It invokes the **RunTask** API on the crema-ecs-cluster.
  - **Override:** It loads the crema-rake-tasks definition but **replaces** the default echo command with ["bundle", "exec", "rake", "subscription:renew"].
  - **Execution:** Fargate provisions a container, runs the specific rake task, and streams output to CloudWatch Logs (/ecs/crema-rake-tasks/rake/...).
  - **Completion:** The task finishes successfully and the container terminates automatically (saving costs compared to a constantly running worker).

## Phase 5: Data Migration (Big Bang Strategy)

**Objective:** Migrate the 10GB PostgreSQL database with data consistency.

- **Assumption:** A "Big Bang" approach is approved. This implies a scheduled maintenance window where the site is read-only or offline.
1. **Preparation:**
    - Notify stakeholders of the maintenance window.
  2. **Execution (Maintenance Window):**
    - **Stop Traffic:** Enable maintenance mode on Heroku to stop writes.
    - **Export:** Create a fresh dump of the Heroku Postgres database.
    - **Import:** Restore the dump to the new Amazon RDS instance using pg\_restore.
    - **Validation:** Verify row counts on critical tables (Users, Orders, Subscriptions).

## Phase 6: Redis Strategy

**Objective:** Stand up the cache layer.

- **Strategy:** Fresh Start.
- **Assumption:** As confirmed, Redis is used strictly as a cache/session store. No data migration is required.
- **Action:** Ensure the Rails config/cable.yml and config/redis.yml (or ENV vars) are updated to point to the new ElastiCache endpoint. Old sessions will be invalidated, requiring users to log in again post-migration.

## Phase 7: Deployment & Cutover

**Objective:** Go live on AWS.

1. **Deploy Application:**
  - Trigger the manual release or push to master to initiate the new CI/CD pipeline.
  - Wait for Health Checks to pass on the Target Group.
2. **Testing:**

- Access the application via the ALB DNS name.
  - Verify connectivity to RDS and Redis.
  - Test critical flows: Login, View Subscription, Checkout (using test mode if possible, or verify Stripe API connectivity).
3. **DNS Cutover:**
- Update Route53 (or current DNS provider) to point the domain to the AWS ALB.
  - Provision SSL Certificates via ACM (AWS Certificate Manager) attached to the ALB.